

How compiler driven technologies can be useful to speedup data processing in python

Sep 20, 2024

Sourav Saha (NEC)

Agenda

- ◆ Icebreaking
- ◆ About Pandas
- ◆ Tips and Tricks of Optimizing Large-scale Data processing workload
- ◆ Compiler driven technologies to optimize the problems
- ◆ FireDucks and Its Offerings
- ◆ FireDucks Optimization Strategy
- ◆ Evaluation Benchmarks
- ◆ Resources on FireDucks
- ◆ Test Drive
- ◆ FAQs

Quick Introduction!



SOURAV SAHA – Research Engineer @ NEC Corporation

<https://www.linkedin.com/in/sourav-%E3%82%BD%E3%82%A6%E3%83%A9%E3%83%96-saha-%E3%82%B5%E3%83%8F-a5750259/>

<https://twitter.com/SouravSaha97589>

Hello, I am a software professional with 11+ years of working experience across diverse areas of **HPC, Vector Supercomputing, Distributed Programming, Big Data and Machine Learning**. Currently, my team at NEC R&D Lab, Japan, is researching various data processing-related algorithms. Blending the mixture of different niche technologies related to compiler framework, high-performance computing, and multi-threaded programming, we have developed a Python library named FireDucks with highly compatible pandas APIs for DataFrame-related operations.



Mr. Kazuhisa Ishizaka
(Primary Author)

we wanted to
develop some library
using compiler
technology

we wanted to
speed-up python

Data
Scientists
often face
issues with
slow
performance
of pandas



User Program

pandas API

FireDucks

groupby

join

dropna

filter

sort

corr

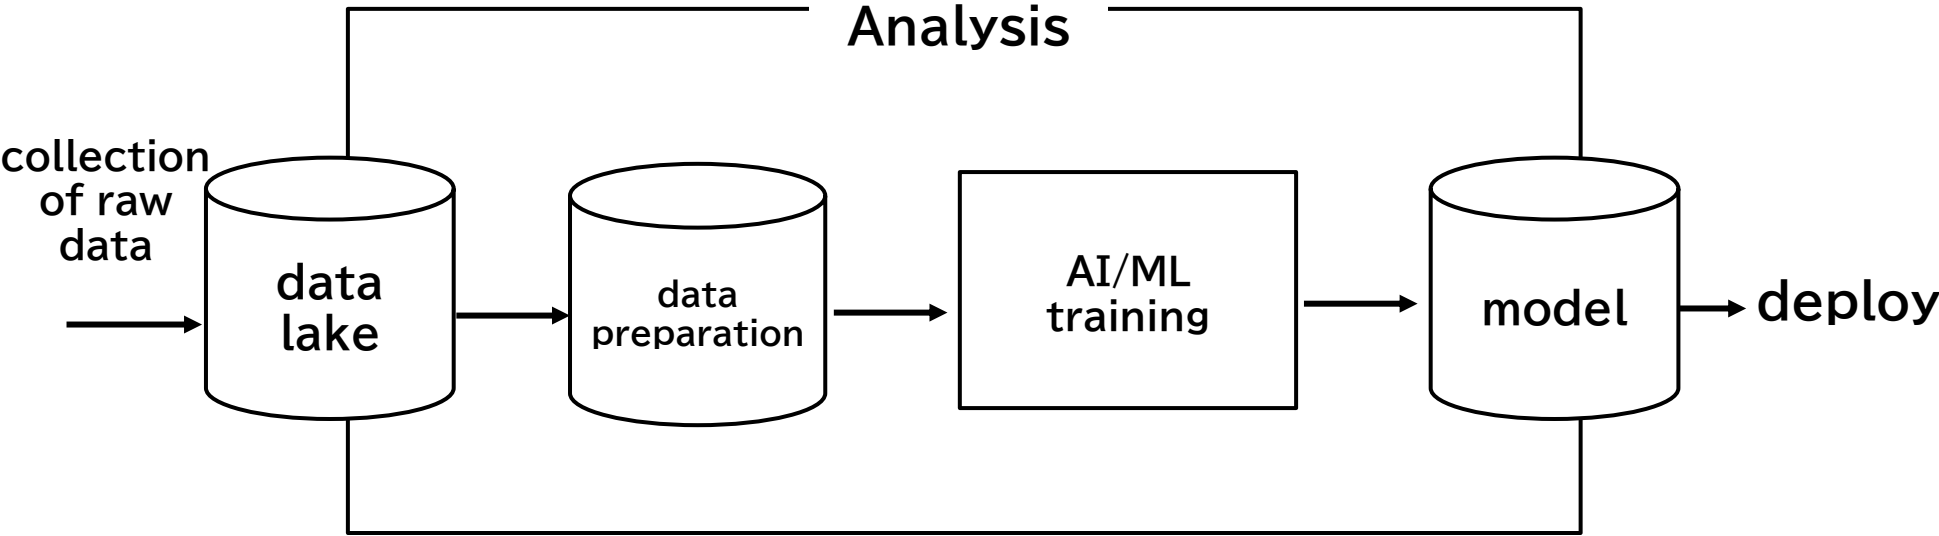

compiler
technologies



<https://www.nec.com/en/global/solutions/hpc/sx/index.html>

Workflow of a Data Scientist


almost 75% efforts of a Data Scientist spent on data preparation



Anaconda:
The State of Data Science 2020

About Pandas (1/2)

◆ Most popular Python library for data analytics.




- It (mostly) doesn't support parallel computation.
- It doesn't have any auto-optimization feature.
- Hence, it is not suitable for processing large datasets.
- Very slow execution reduces the efficiency of a data analyst.
- Long-running execution
 - produces higher cloud costs
 - attributes to higher CO2 emission



About Pandas (2/2)

◆ Most popular Python library for data analytics.



The way of implementing a query in pandas-like library (that does not support query optimization) heavily impacts its performance!!



- We will discuss a couple of approaches to improve the performance related to computational time and memory of a query written in pandas, when processing large-scale data.
- We will also discuss how those approaches can be automated using compiler technologies.

Ice-Breaking Session

(test your pandas skill)

Quick check on basic pandas operations (1/5)

◆ Which one of the following is the right method of getting top-2 rows based on the column “A” from table “df”?

1. `df.sort("A", ascending=True).head(2)`
2. `df["A"].top_k(2)`
3. `df.sort("A", ascending=False).first(2)`
4. `df.sort_values("A", ascending=False).head(2)`

	A	B
0	2	10
1	5	30
2	1	20
3	3	70
4	7	60
5	8	40
6	4	80




	A	B
5	8	40
4	7	60

Quick check on basic pandas operations (2/5)

- ◆ Which ones of the following are the right methods of performing inner-join of table “df1” with table “df2” on common key-column “x”?

1. `pd.merge(df1, df2, on="x", how="inner")`
2. `df1.inner_join(df2, on="x")`
3. `df1.merge(df2, on="x", how="inner")`
4. `df1.merge(df2, on="x")`



Quick check on basic pandas operations (3/5)

◆ Which one of the following is the right method to remove rows having a missing value?

1. `df.dropna()`
2. `df.dropna(how="any")`
3. `df[~df["A"].isnull()]`
4. All of the above

	A	B
0	N	10
1	5	30
2	N	20
3	3	70
4	7	60
5	8	40
6	4	80



	A	B
1	5	30
3	3	70
4	7	60
5	8	40
6	4	80

Quick check on basic pandas operations (4/5)

◆ Which one of the following is the right method of selecting columns “A”, “D” and “E” from table “df”?

1. `df[["A", "D", "E"]]`
2. `df.loc[:, ["A", "D", "E"]]`
3. `df.iloc[:, [0, 3, 4]]`
4. All of the above

	A	B	C	D	E
0	2	10	10	g	9
1	5	30	69	a	2
2	1	20	31	g	8
3	3	70	45	f	3
4	7	60	59	e	1
5	8	40	66	f	1
6	4	80	97	h	8



	A	D	E
0	2	g	9
1	5	a	2
2	1	g	8
3	3	f	3
4	7	e	1
5	8	f	1
6	4	h	8

Quick check on basic pandas operations (5/5)

◆ Select the options for appending a new column "F" by doubling the column "B" from table "df".

1. `df["F"] = df["B"] * 2`
2. `df.assign(F=lambda x: x["B"] * 2)`
3. `df.with_columns(df.col("B") * 2).alias("F")`
4. `df.insert(5, "F", df["B"]*2)`

	A	B	C	D	E
0	2	10	10	g	9
1	5	30	69	a	2
2	1	20	31	g	8
3	3	70	45	f	3
4	7	60	59	e	1
5	8	40	66	f	1
6	4	80	97	h	8



	A	B	C	D	E	F
0	2	10	10	g	9	20
1	5	30	69	a	2	60
2	1	20	31	g	8	40
3	3	70	45	f	3	140
4	7	60	59	e	1	120
5	8	40	66	f	1	80
6	4	80	97	h	8	160

Performance Challenges & Best Practices to follow

(1) importance of chained expression

```
def foo(filename):  
    df = pd.read_csv(filename)  
    t1 = df.drop_duplicates()  
    t2 = t1.sort_values("B")  
    t3 = t2.head(2)  
    return t3
```



re-write using chained
expression

```
def foo(filename):  
    return (  
        pd.read_csv(filename)  
        .drop_duplicates()  
        .sort_values("B")  
        .head(2)  
    )
```

df: ~16 GB

A	B	C
u	0.91	1
a	1.00	4
a	1.00	4
o	0.24	0
o	0.24	0
e	0.43	1
u	0.91	1
e	0.20	2
o	0.24	0
a	1.00	4

t1: ~8 GB

A	B	C
u	0.91	1
a	1.00	4
o	0.24	0
e	0.43	1
e	0.20	2

t3: ~8 GB

A	B	C
a	1.00	4
u	0.91	1
e	0.43	1
o	0.24	0
e	0.20	2

t4: ~x KB

A	B	C
a	1.00	4
u	0.91	1

drop_duplicates

sort

head(2)

A	B	C
u	0.91	1
a	1.00	4
a	1.00	4
o	0.24	0
o	0.24	0
e	0.43	1
u	0.91	1
e	0.20	2
o	0.24	0
a	1.00	4

A	B	C
u	0.91	1
a	1.00	4
o	0.24	0
e	0.43	1
e	0.20	2

A	B	C
a	1.00	4
u	0.91	1
e	0.43	1
o	0.24	0
e	0.20	2

A	B	C
a	1.00	4
u	0.91	1

drop_duplicates

sort

head(2)

challenges with pandas APIs when writing chained expression

```
def foo(filename):  
    df = pd.read_csv(filename)  
    t1 = df.drop_duplicates()  
    t2 = t1[t1["B"] > 0.20]  
    t3 = t2.sort_values("B")  
    t4 = t3.head(2)  
    return t4
```



re-write using chained expression

```
def foo(filename):  
    return (  
        pd.read_csv(filename)  
        .drop_duplicates()  
        .??  
        .sort_values("B")  
        .head(2)  
    )
```

df: ~16 GB

A	B	C
u	0.91	1
a	1.00	4
a	1.00	4
o	0.24	0
o	0.24	0
e	0.43	1
u	0.91	1
e	0.20	2
o	0.24	0
a	1.00	4

t1: ~8 GB

A	B	C
u	0.91	1
a	1.00	4
o	0.24	0
e	0.43	1
e	0.20	2

t2: ~8 GB

A	B	C
u	0.91	1
a	1.00	4
o	0.24	0
e	0.43	1

t3: ~8 GB

A	B	C
a	1.00	4
u	0.91	1
e	0.43	1
o	0.24	0

t4: ~x KB

A	B	C
a	1.00	4
u	0.91	1

drop_duplicates

filter

sort

head(2)

```
def foo(filename):  
    return (  
        pd.read_csv(filename)  
        .drop_duplicates()  
        .query("B > 0.20")  
        .sort_values("B")  
        .head(2)  
    )
```

```
def foo(filename):  
    return (  
        pd.read_csv(filename)  
        .drop_duplicates()  
        .pipe(lambda tmp: tmp[tmp["B"] > 0.20])  
        .sort_values("B")  
        .head(2)  
    )
```

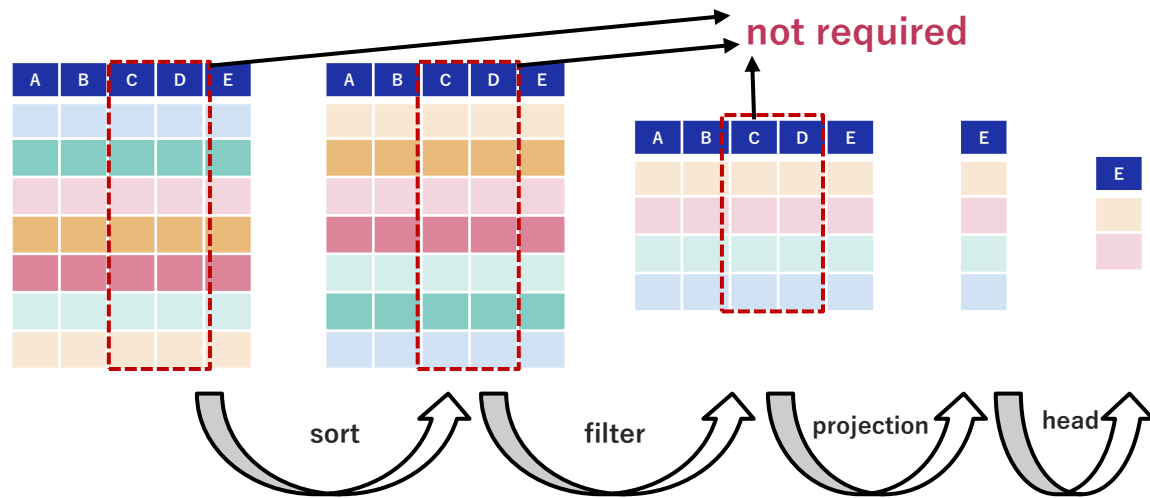
query(): allows you to write SQL-like conditional expression, helping you to perform filter on the current state of the input frame, but its a little slower as it parses the input string to construct the filter mask.

pipe(): a convenient method allowing you to perform a given operation (like filter etc.) on the current state of the input frame without introducing computational overhead.

(2) importance of execution order

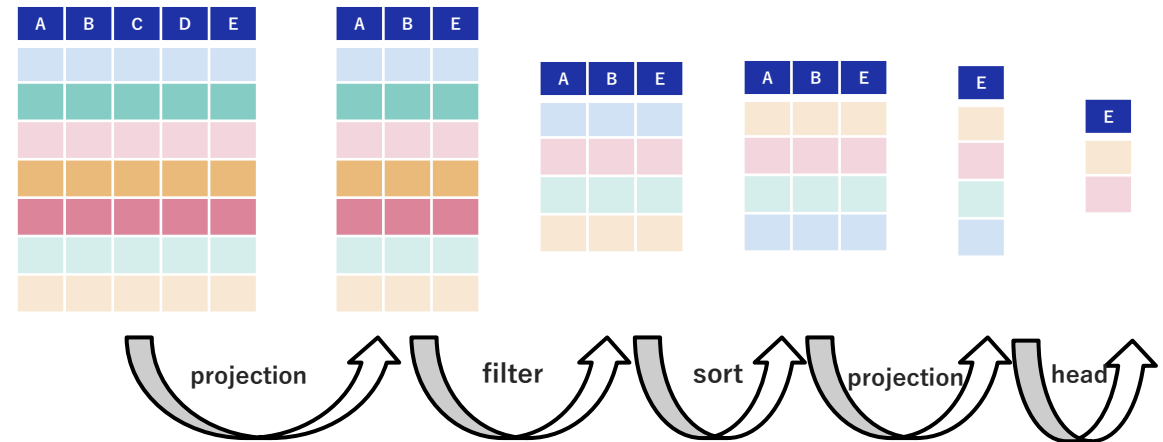
```
df.sort_values("A")  
.query("B > 1")["E"]  
.head(2)
```

※ *sort-order: yellow->red->green->blue*
※ *B=1 for darker shade, B=2 for lighter shade*



SAMPLE QUERY

```
df.loc[:, ["A", "B", "E"]]  
.query("B > 1")  
.sort_values("A")["E"]  
.head(2)
```



reduction in the
number of columns
(**projection
pushdown**)

reduction in the
number of rows
(**predicate
pushdown**)

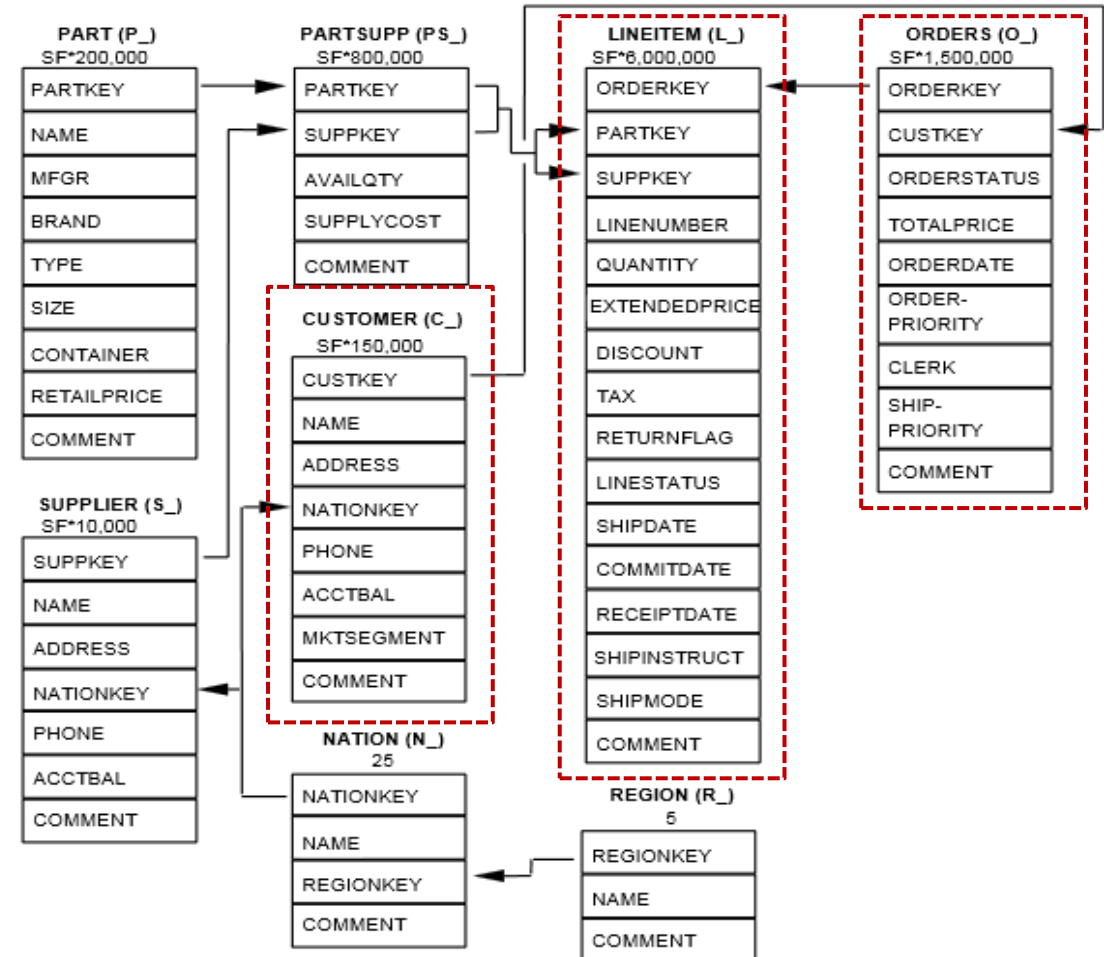
**OPTIMIZED
QUERY**

Exercise: Query #3 from TPC-H Benchmark (SQL -> pandas)

- ◆ [query to retrieve the 10 unshipped orders with the highest value.](#)

```
SELECT l_orderkey,  
       sum(l_extendedprice * (1 - l_discount)) as revenue,  
       o_orderdate,  
       o_shippriority  
FROM customer, orders, lineitem  
WHERE  c_mktsegment = 'BUILDING' AND  
       c_custkey = o_custkey AND  
       l_orderkey = o_orderkey AND  
       o_orderdate < date '1995-03-15' AND  
       l_shipdate > date '1995-03-15'  
GROUP BY l_orderkey, o_orderdate, o_shippriority  
ORDER BY revenue desc, o_orderdate  
LIMIT 10;
```

```
rescols = ["l_orderkey", "revenue", "o_orderdate", "o_shippriority"]  
result = (  
    customer.merge(orders, left_on="c_custkey", right_on="o_custkey")  
    .merge(lineitem, left_on="o_orderkey", right_on="l_orderkey")  
    .pipe(lambda df: df[df["c_mktsegment"] == "BUILDING"])  
    .pipe(lambda df: df[df["o_orderdate"] < datetime(1995, 3, 15)])  
    .pipe(lambda df: df[df["l_shipdate"] > datetime(1995, 3, 15)])  
    .assign(revenue=lambda df: df["l_extendedprice"] * (1 - df["l_discount"]))  
    .groupby(["l_orderkey", "o_orderdate", "o_shippriority"], as_index=False)  
    .agg({"revenue": "sum"})[rescols]  
    .sort_values(["revenue", "o_orderdate"], ascending=[False, True])  
    .head(10)  
)
```



Exercise: Query #3 from TPC-H Benchmark (pandas -> optimized pandas)

```
rescols = ["l_orderkey", "revenue", "o_orderdate", "o_shippriority"]
result = (
    customer.merge(orders, left_on="c_custkey", right_on="o_custkey")
    .merge(lineitem, left_on="o_orderkey", right_on="l_orderkey")
    .pipe(lambda df: df[df["c_mktsegment"] == "BUILDING"])
    .pipe(lambda df: df[df["o_orderdate"] < datetime(1995, 3, 15)])
    .pipe(lambda df: df[df["l_shipdate"] > datetime(1995, 3, 15)])
    .assign(revenue=lambda df: df["l_extendedprice"] * (1 - df["l_discount"]))
    .groupby(["l_orderkey", "o_orderdate", "o_shippriority"], as_index=False)
    .agg({"revenue": "sum"})[rescols]
    .sort_values(["revenue", "o_orderdate"], ascending=[False, True])
    .head(10)
)
```

Exec-time: 68.55 s

Scale Factor: 10

6.5x

Exec-time: 10.33 s

```
# projection-filter: to reduce scope of "customer" table to be processed
cust = customer[["c_custkey", "c_mktsegment"]] # (2/8)
f_cust = cust[cust["c_mktsegment"] == "BUILDING"]

# projection-filter: to reduce scope of "orders" table to be processed
ord = orders[["o_custkey", "o_orderkey", "o_orderdate", "o_shippriority"]] (4/9)
f_ord = ord[ord["o_orderdate"] < datetime(1995, 3, 15)]

# projection-filter: to reduce scope of "lineitem" table to be processed
litem = lineitem[["l_orderkey", "l_shipdate", "l_extendedprice", "l_discount"]] (4/16)
f_litem = litem[litem["l_shipdate"] > datetime(1995, 3, 15)]

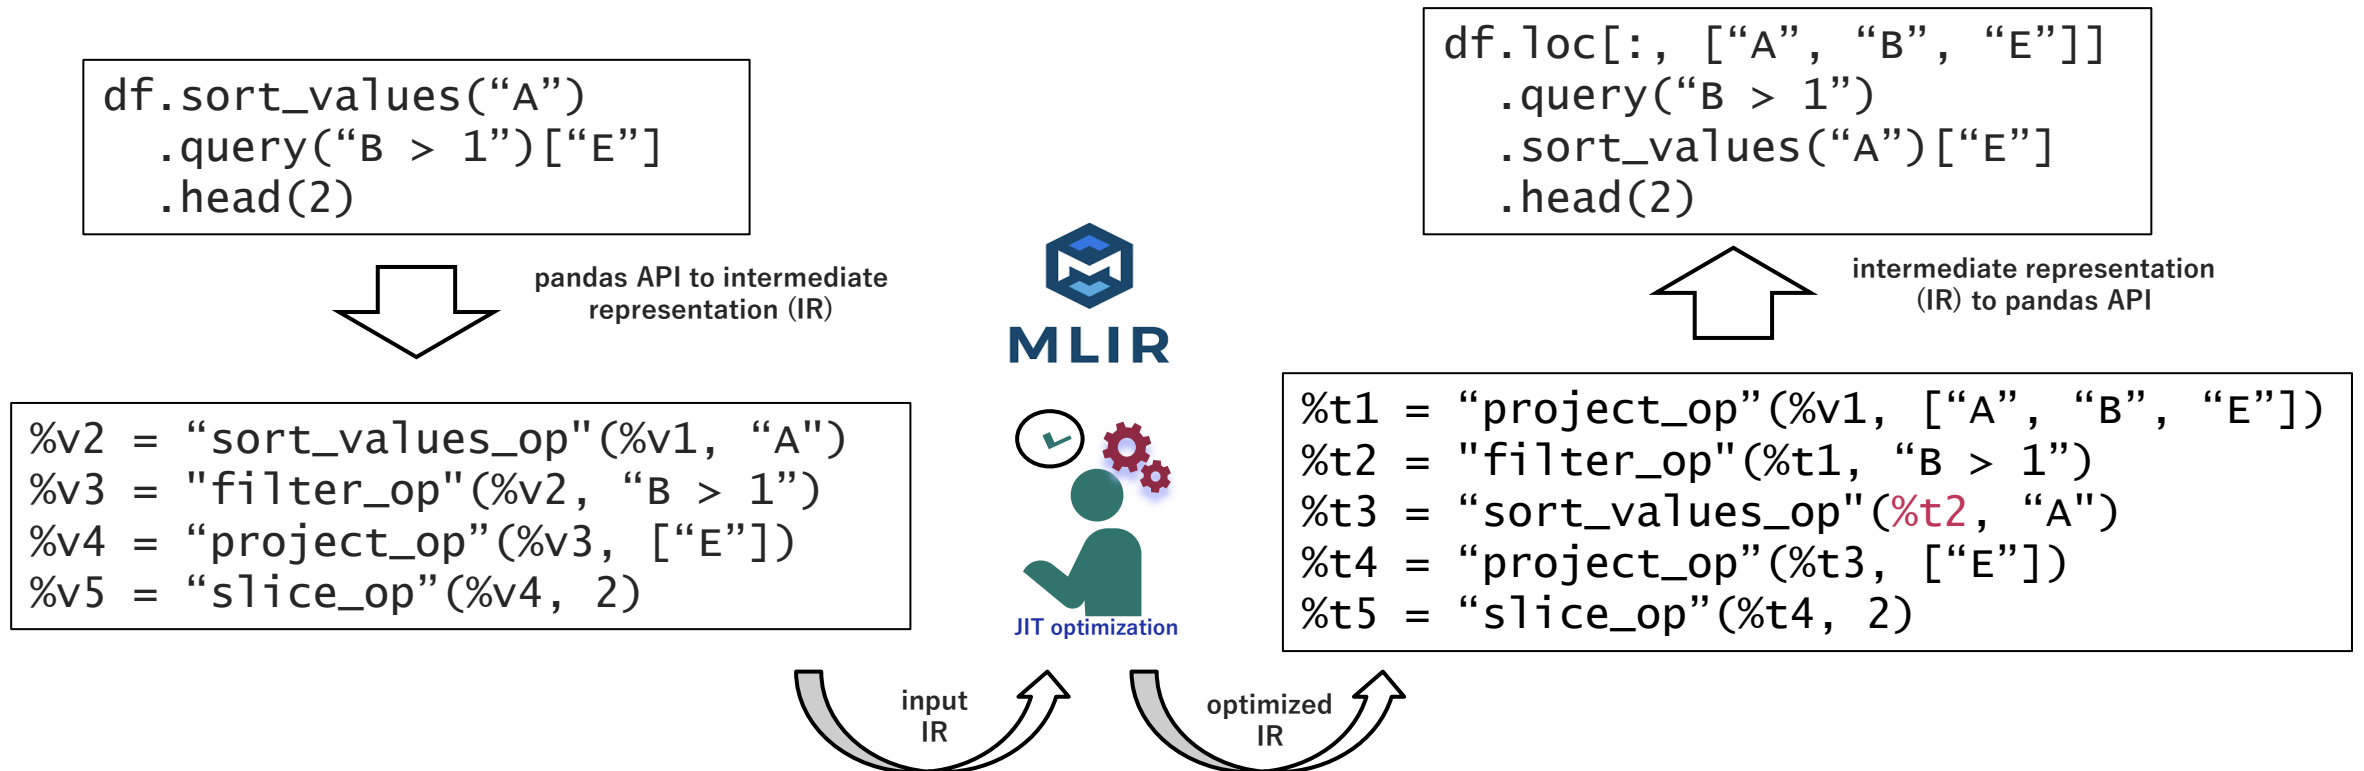
rescols = ["l_orderkey", "revenue", "o_orderdate", "o_shippriority"]
result = ( f_cust.merge(f_ord, left_on="c_custkey", right_on="o_custkey")
    .merge(f_litem, left_on="o_orderkey", right_on="l_orderkey")
    .assign(revenue=lambda df: df["l_extendedprice"] * (1 - df["l_discount"]))
    .pipe(lambda df: df[rescols])
    .groupby(["l_orderkey", "o_orderdate", "o_shippriority"], as_index=False)
    .agg({"revenue": "sum"})[rescols]
    .sort_values(["revenue", "o_orderdate"], ascending=[False, True])
    .head(10)
)
```

Automatic Optimization

Idea #1

- **Can such optimization be automated?**


- Yes, using LLVM/MLIR define-by-run mechanism we can build specialized intermediate representation for each pandas API.
- The generated IRs can be parsed to implement different domain-specific optimizations, such as projection pushdown, predicate pushdown, etc.
- the optimized IRs can be translated back to the pandas API.



Idea #2

- Pandas methods are slow due to poor memory utilization and single-core computation.
- But pandas is one of the most popular data manipulation tools.
- **How can we solve the core performance issue in pandas while keeping the same API for users?**
 - Well, we can
 - have a frontend with pandas API that generates IR.
 - develop our own library parallelizing the workload of DataFrame-related methods as a backend.
 - translate the optimized IRs to the **backend library API** (instead of pandas API).

```
df.sort_values("A")  
.query("B > 1")["E"]  
.head(2)
```




frontend pandas API to
intermediate representation
(IR)

```
%v2 = "sort_values_op"(%v1, "A")  
%v3 = "filter_op"(%v2, "B > 1")  
%v4 = "project_op"(%v3, ["E"])  
%v5 = "slice_op"(%v4, 2)
```



JIT optimization



```
t1 = backend::project_columns(df, {"A", "B", "C"});  
t2 = backend::filter_rows(t1, "B > 1");  
t3 = backend::sort_values(t2, "A");  
t4 = backend::project_columns(t3, {"E"});  
t5 = backend::slice_rows(t4, 2);
```



intermediate representation
(IR) to backend API

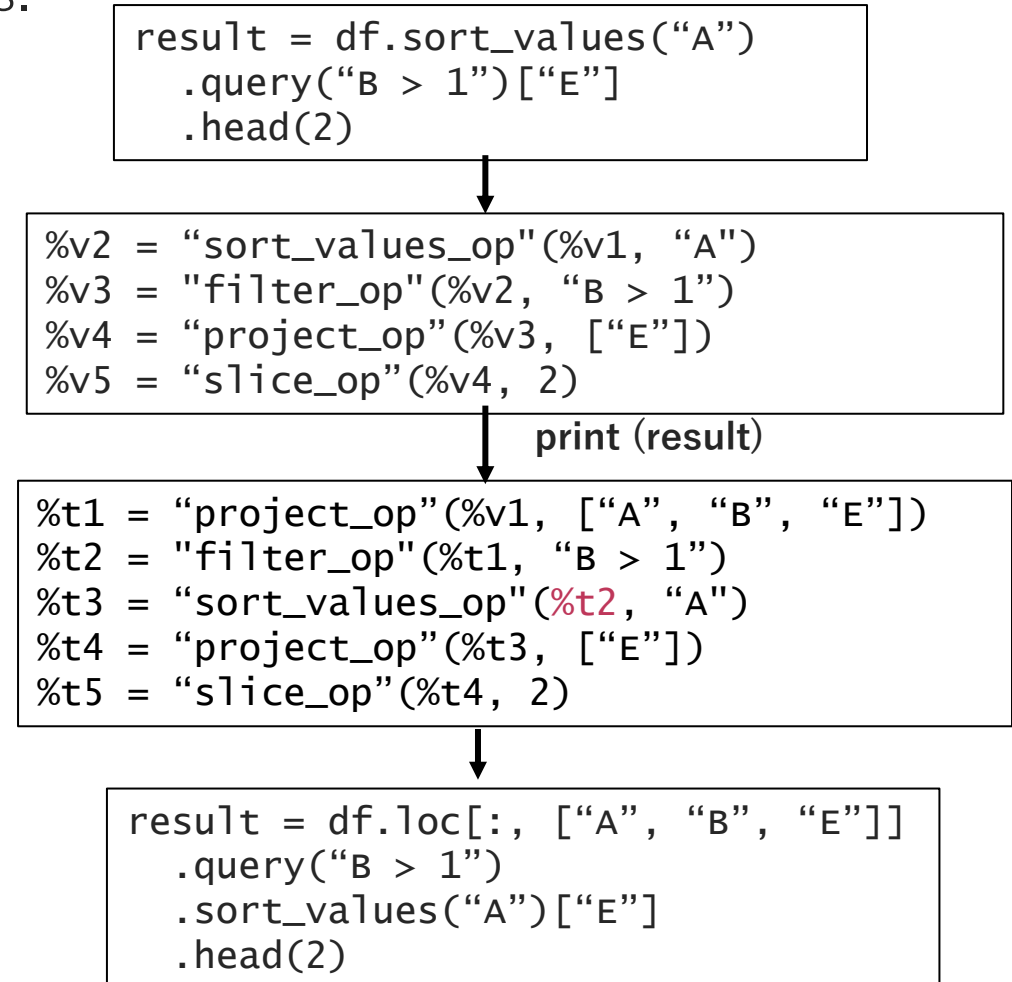
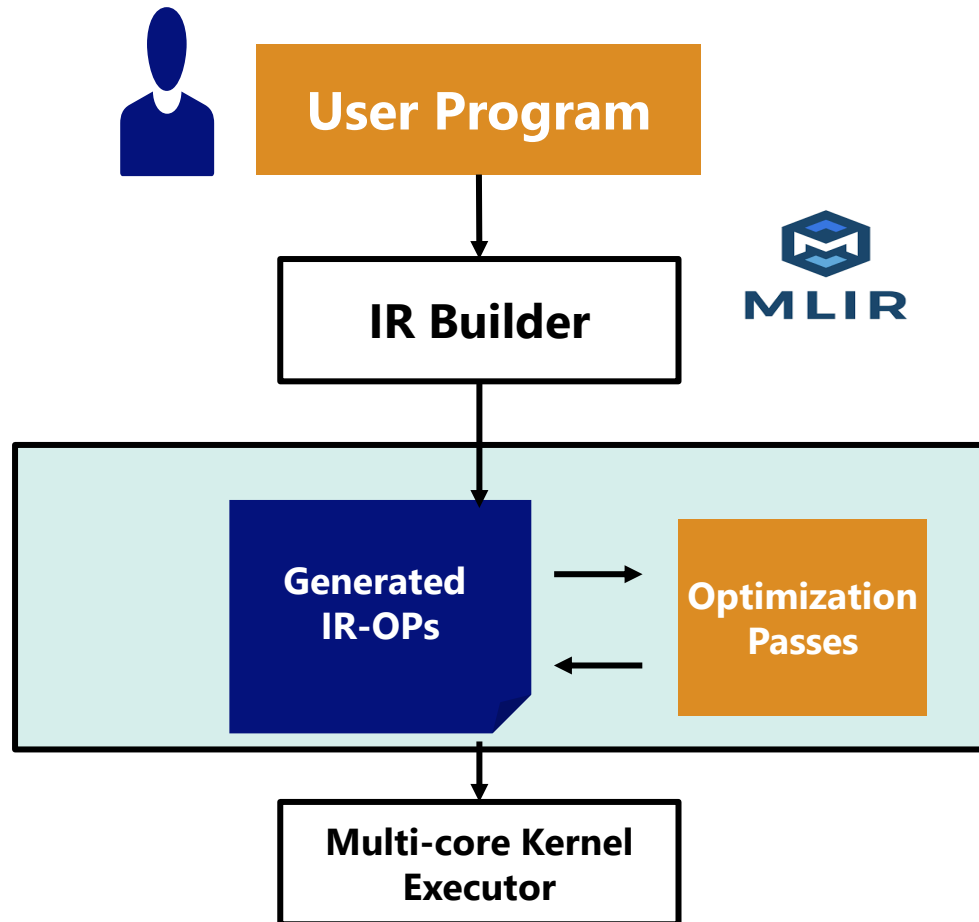
```
%t1 = "project_op"(%v1, ["A", "B", "E"])  
%t2 = "filter_op"(%t1, "B > 1")  
%t3 = "sort_values_op"(%t2, "A")  
%t4 = "project_op"(%t3, ["E"])  
%t5 = "slice_op"(%t4, 2)
```

Introducing FireDucks

Introducing FireDucks

※IR: Intermediate Representation

FireDucks (Flexible **IR** Engine for DataFrame) is a high-performance compiler-accelerated DataFrame library with highly compatible pandas APIs.



Primary Objective: Write Once, Execute Anywhere


Why FireDucks?

※IR: Intermediate Representation

FireDucks (Flexible **IR** Engine for DataFrame) is a high-performance compiler-accelerated DataFrame library with highly compatible pandas APIs.

Speed: significantly faster than pandas

- FireDucks is multithreaded to fully exploit the modern processor
- Lazy execution model with Just-In-Time optimization using a defined-by-run mechanism supported by MLIR (a subproject of LLVM).
 - supports both lazy and non-lazy execution models without modifying user programs (same API).



Ease of use: drop-in replacement of pandas

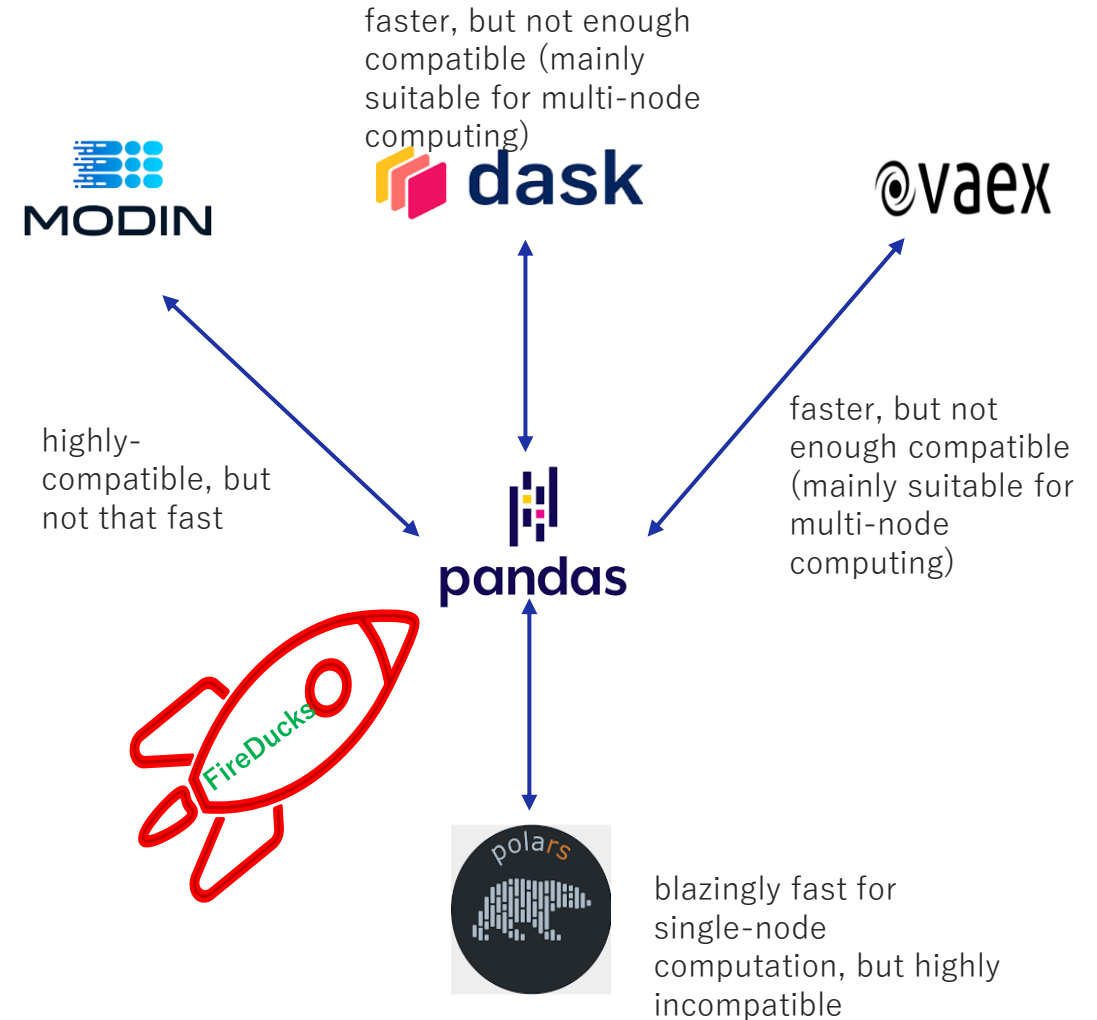
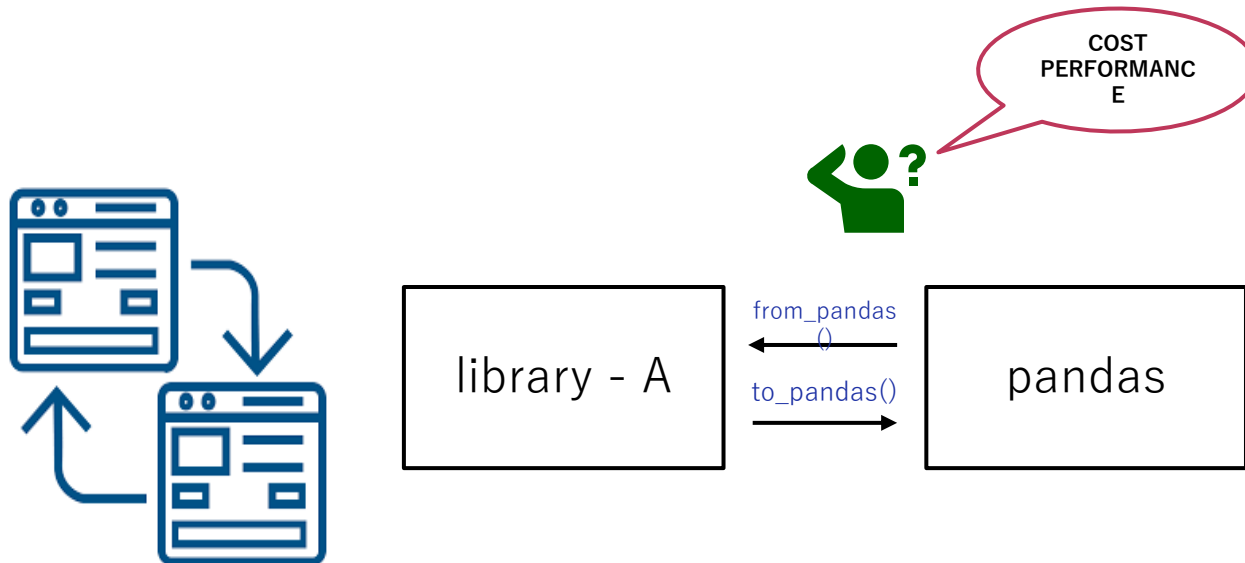
- FireDucks is highly compatible with pandas API
 - seamless integration is possible not only for an existing pandas program but also for any external libraries (like seaborn, scikit-learn, etc.) that internally use pandas dataframes.
- No extra learning is required
- No code modification is required



Seamless Integration with pandas

Three most common challenges in switching from pandas:

- Needs to learn new library and their interfaces.
- Manual fallback to pandas when the target library doesn't support a method used in an existing pandas application.
- Performance can be evaluated, and results can be tested after the migration is completed.

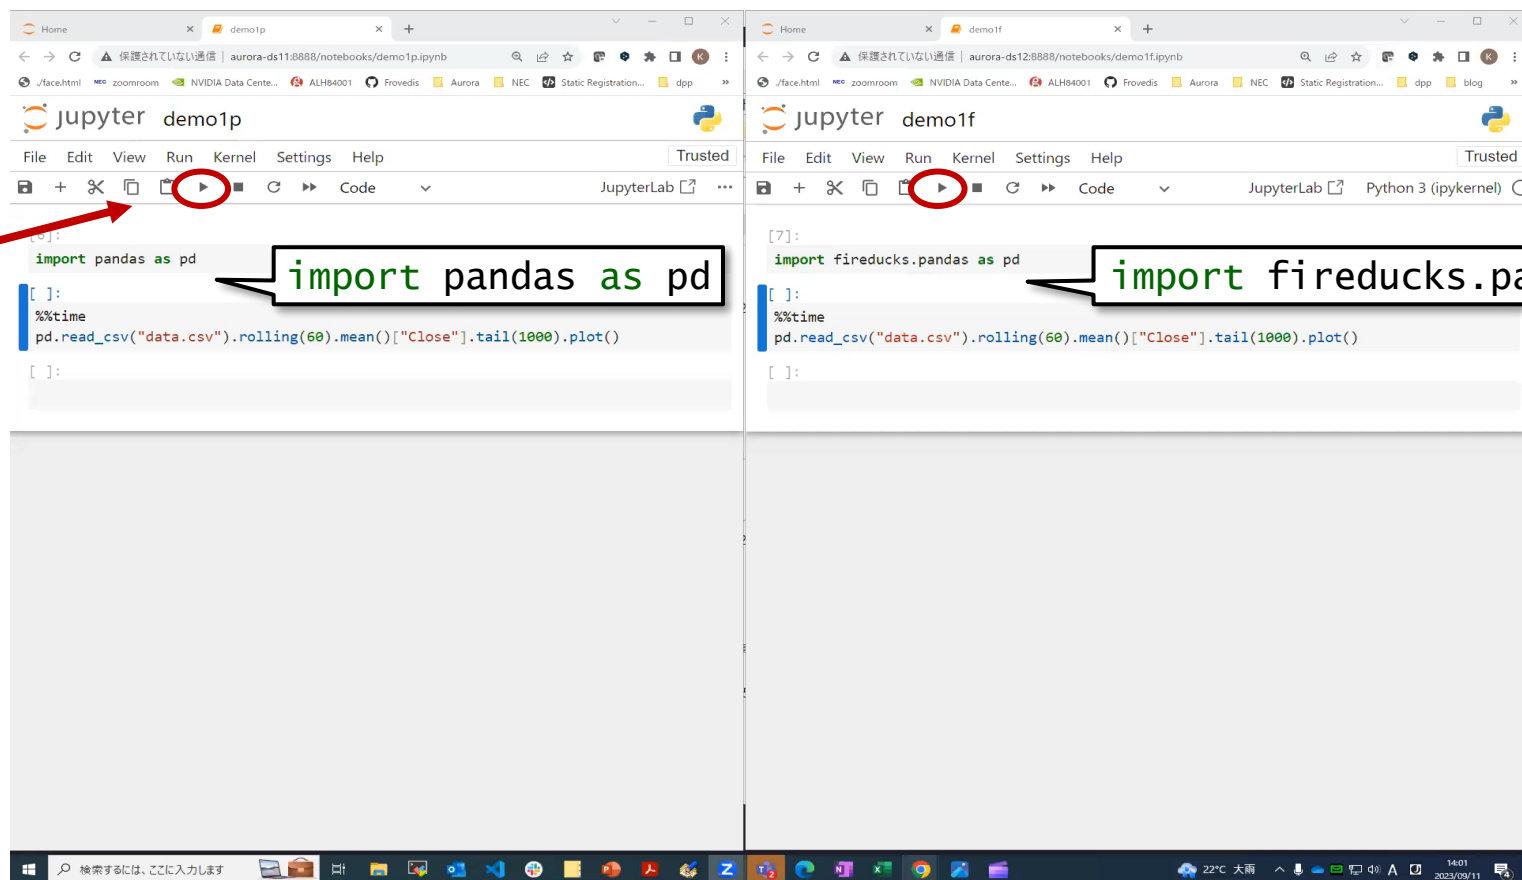


Let's Have a Quick Demo!

```
pd.read_csv("data.csv").rolling(60).mean()["Close"].tail(1000).plot()
```

pandas the difference is only in the import **FireDucks**

button to
start
execution



Program to
calculate
moving average

pandas: 4.06s

↓ ~15x

FireDucks: 275ms

data.csv:
[Bitcoin Historical Data](#)

Usage of FireDucks

1. Explicit Import

easy to import

```
# import pandas as pd
import fireducks.pandas as pd
```

simply change the import statement

2. Import Hook

FireDucks provides command line option to automatically replace "**pandas**" with "**fireducks.pandas**"

```
$ python -m fireducks.pandas program.py
```

zero code modification

```
import mod_A
import mod_B
import mod_C
import pandas as
pd
:
```

program.py

```
import pandas as
pd
:
import pandas as
pd
:
import pandas as
pd
:
```

mod_A.py

mod_B.py

mod_C.py

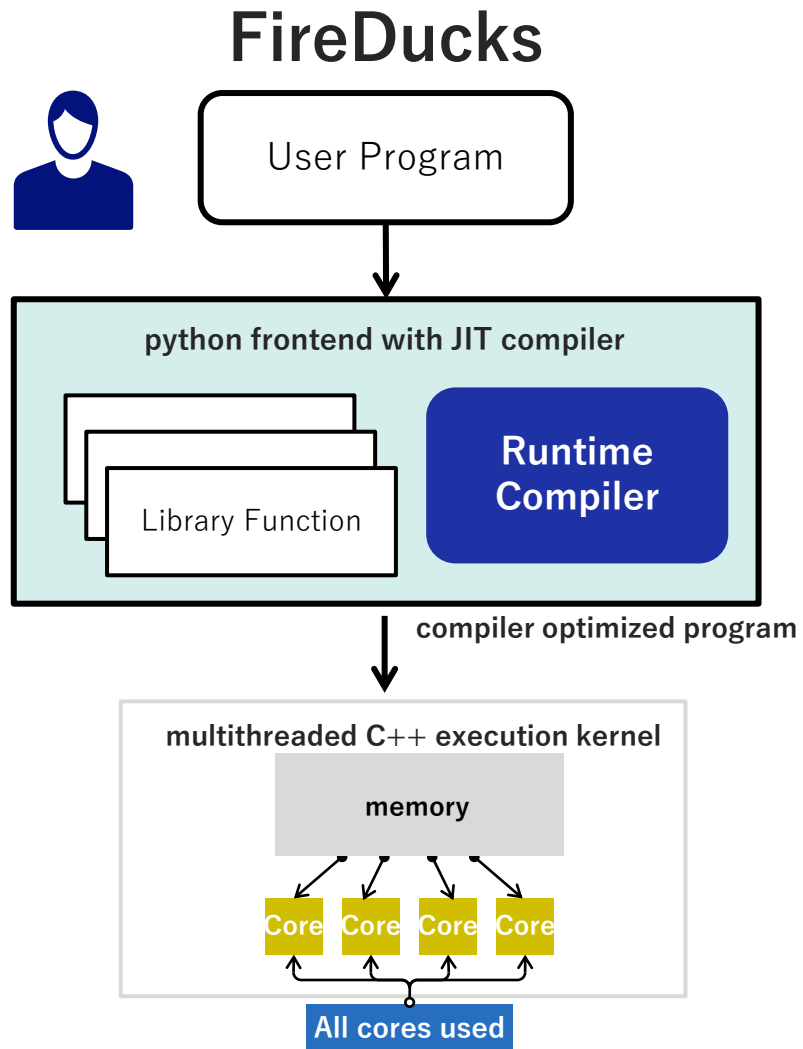
3. Notebook Extension

FireDucks provides simple import extension for interactive notebooks.

```
%load_ext fireducks.pandas
import pandas as pd
```

simple integration in a notebook

Optimization Features



1. **Compiler Specific Optimizations:** Common Sub-expression Elimination, Dead-code Elimination, Constant Folding etc.
2. **Domain Specific Optimization:** Optimization at query-level: reordering instructions etc.
3. **Pandas Specific Optimization:** selection of suitable pandas APIs, selection of suitable parameter etc.

1. **Multi-threaded Computation:** Leverage all the available computational cores.
2. **Efficient Memory Management:** Data Structures backed by Apache Arrow
3. **Optimized Kernels:** Patented algorithms for Database like kernel operations: like sorting, join, filter, groupby, dropna etc. developed in C++ from scratch.

IR-driven Lazy-execution addresses memory issue with intermediate tables

```
def foo(filename):  
    df = pd.read_csv(filename)  
    t1 = df.drop_duplicates()  
    t2 = t1[t1["B"] > 0.20]  
    t3 = t2.sort_values("B")  
    t4 = t3.head(2)  
    return t4
```

```
ret = foo("data.csv")  
print(ret.shape)
```

example without chained
expression

```
def foo(filename):  
    return (  
        pd.read_csv(filename)  
        .drop_duplicates()  
        .query("B > 0.20")  
        .sort_values("B")  
        .head(2)  
    )
```

```
ret = foo("data.csv")  
print(ret.shape)
```

example with chained expression

```
%t3 = read_csv_with_metadata('dummy.csv', ...)  
%t4 = drop_duplicates(%t3, ...)  
%t5 = project(%t4, 'B')  
%t6 = gt.vector.scalar(%t5, 0.20)  
%t7 = filter(%t4, %t6)  
%t8 = sort_values(%t7, ['B'], [True])  
%t9 = slice(%t8, 0, 2, 1)  
%v10 = get_shape(%t9)  
return(%t9, %v10)
```

IR Generated by FireDucks

(can be inspected when setting environment variable FIRE_LOG_LEVEL=3)

Compiler Specific Optimizations

- same operation on the same data repeatedly

Find year and month-wise average sales

```
s = pd.Series(["2020-01-01", "2021-01-01", "2022-01-01"])
```

```
df = pd.DataFrame()
```

```
df["year"] = pd.to_datetime(s).dt.year
```

```
df["month"] = pd.to_datetime(s).dt.month
```

```
df["sales"] = [100, 200, 500]
```

```
r = df.groupby(["year", "month"])["sales"].mean()
```

```
print(r)
```

```
%t8 = to_datetime(%t7, None)
```

```
%t9 = datetime_extract(%t8, 'year')
```

```
%t10 = setitem(%t6, 'year', %t9)
```

```
%t11 = to_datetime(%t7, None)
```

```
%t12 = datetime_extract(%t11, 'month')
```

```
%t13 = setitem(%t10, 'month', %t12)
```

```
%t14 = from_pandas.frame.metadata(%arg4, %arg5)
```

```
%t15 = setitem(%t13, 'sales', %t14)
```

```
%t16 = groupby_select_agg(%t15, ['year', 'month'], ['mean'], [], [], 'sales')
```

```
%v17 = get_shape(%t16)
```

```
return(%t16, %v17)
```

Common Sub-expression Elimination

Find year and month-wise average sales

```
s = pd.Series(["2020-01-01", "2021-01-01", "2022-01-01"])
```

```
tmp = pd.to_datetime(s).
```

```
df = pd.DataFrame()
```

```
df["year"] = tmp.dt.year
```

```
df["month"] = tmp.dt.month
```

```
df["sales"] = [100, 200, 500]
```

```
r = df.groupby(["year", "month"])["sales"].mean()
```

```
print(r)
```

```
%t8 = to_datetime(%t7, None)
```

```
%t9 = datetime_extract(%t8, 'year')
```

```
%t11 = setitem(%t6, 'year', %t9)
```

```
%t12 = datetime_extract(%t8, 'month')
```

```
%t14 = setitem(%t11, 'month', %t12)
```

```
%t15 = from_pandas.frame.metadata(%arg4, %arg5)
```

```
%t16 = project(%t14, ['year', 'month'])
```

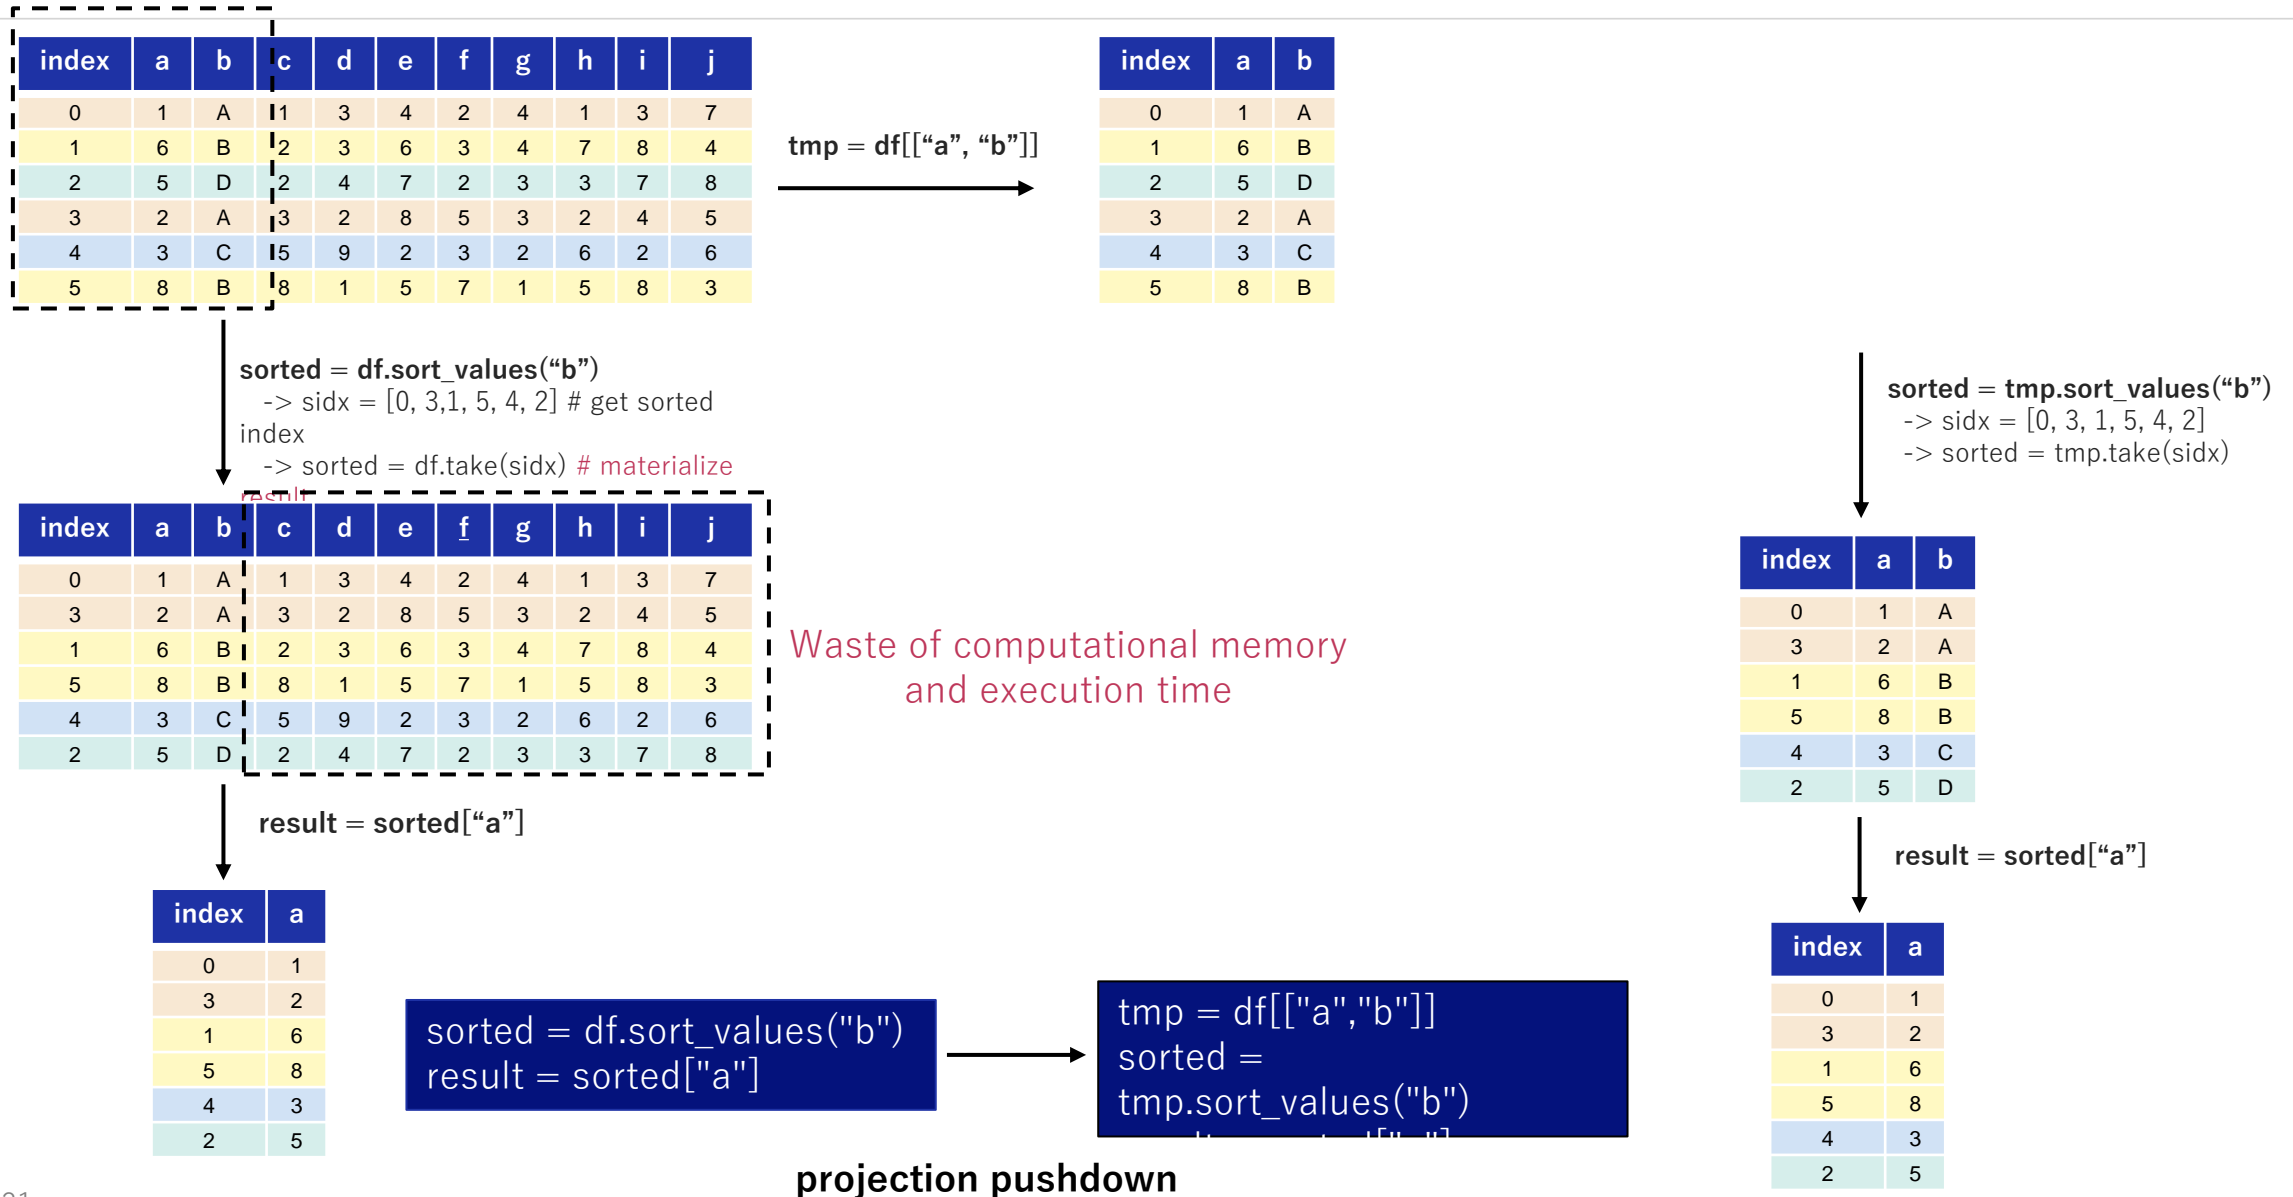
```
%t17 = setitem(%t16, 'sales', %t15)
```

```
%t18 = groupby_select_agg(%t17, ['year', 'month'], ['mean'], [], [], 'sales')
```

```
%v19 = get_shape(%t18)
```

```
return(%t18, %v19)
```

Domain Specific Optimization (Example #1)



Domain Specific Optimization (Example #2) (1/2)

ID	E_Name	Gender	C_Code
1	A	Male	1
2	B	Male	1
3	C	Female	2
4	E	Male	2
5	F	Female	1
6	G	Female	2
7	H	Male	1
8	I	Female	2

employee

C_Code	C_Name
1	India
2	Japan

country

merge

ID	E_Name	Gender	C_Code	C_Name
1	A	Male	1	India
2	B	Male	1	India
3	C	Female	2	Japan
4	E	Male	2	Japan
5	F	Female	1	India
6	G	Female	2	Japan
7	H	Male	1	India
8	I	Female	2	Japan

filter

ID	E_Name	Gender	C_Code	C_Name
1	A	Male	1	India
2	B	Male	1	India
4	E	Male	2	Japan
7	H	Male	1	India

groupby-count

C_Name	E_Name
India	3
Japan	2

```
m = employee.merge(country, on="C_Code")
f = m[m["Gender"] == "Male"]
r = f.groupby("C_Name")["E_Name"].count()
print(r)
```

- sample case: **filter after merge operation**
 - merge is an expensive operation, as it involves data copy.
 - performing merge operation on a large dataset and then filtering the output would involve unnecessary costs in data-copy.

Domain Specific Optimization (Example #2) (2/2)

ID	E_Name	Gender	C_Code
1	A	Male	1
2	B	Male	1
3	C	Female	2
4	E	Male	2
5	F	Female	1
6	G	Female	2
7	H	Male	1
8	I	Female	2

employee

C_Code	C_Name
1	India
2	Japan

country

merge

filter


ID	E_Name	Gender	C_Code
1	A	Male	1
2	B	Male	1
4	E	Male	2
7	H	Male	1

ID	Name	Gender	C_Code	C_Name
1	A	Male	1	India
2	B	Male	1	India
4	E	Male	2	Japan
7	H	Male	1	India

**groupby-
count**

C_Name	E_Name
India	3
Japan	2

```
m = employee.merge(country, on="C_Code")
f = m[m["Gender"] == "Male"]
r = f.groupby("C_Name")["E_Name"].count()
print(r)
```



predicate pushdown

```
f = employee[employee["Gender"] == "Male"]
m = f.merge(country, on="C_Code")
r = m.groupby("C_Name")["E_Name"].count()
print(r)
```

Domain Specific Optimization

```
rescols = ["l_orderkey", "revenue", "o_orderdate", "o_shippriority"]
result = (
    customer.merge(orders, left_on="c_custkey", right_on="o_custkey")
    .merge(lineitem, left_on="o_orderkey", right_on="l_orderkey")
    .pipe(lambda df: df[df["c_mktsegment"] == "BUILDING"])
    .pipe(lambda df: df[df["o_orderdate"] < datetime(1995, 3, 15)])
    .pipe(lambda df: df[df["l_shipdate"] > datetime(1995, 3, 15)])
    .assign(revenue=lambda df: df["l_extendedprice"] * (1 - df["l_discount"]))
    .groupby(["l_orderkey", "o_orderdate", "o_shippriority"], as_index=False)
    .agg({"revenue": "sum"})[rescols]
    .sort_values(["revenue", "o_orderdate"], ascending=[False, True])
    .head(10)
)
```

Exec-time: 68.55 s

Scale Factor: 10

6.5x

Exec-time: 10.33 s

projection-filter: to reduce scope of "customer" table to be processed

```
cust = customer[["c_custkey", "c_mktsegment"]] # (2/8)
```

```
f_cust = cust[cust["c_mktsegment"] == "BUILDING"]
```

projection-filter: to reduce scope of "orders" table to be processed

```
ord = orders[["o_custkey", "o_orderkey", "o_orderdate", "o_shippriority"]] (4/9)
```

```
f_ord = ord[ord["o_orderdate"] < datetime(1995, 3, 15)]
```

projection-filter: to reduce scope of "lineitem" table to be processed

```
litem = lineitem[["l_orderkey", "l_shipdate", "l_extendedprice", "l_discount"]] (4/16)
```

```
f_litem = litem[litem["l_shipdate"] > datetime(1995, 3, 15)]
```

```
rescols = ["l_orderkey", "revenue", "o_orderdate", "o_shippriority"]
```

```
result = ( f_cust.merge(f_ord, left_on="c_custkey", right_on="o_custkey")
```

```
    .merge(f_litem, left_on="o_orderkey", right_on="l_orderkey")
```

```
    .assign(revenue=lambda df: df["l_extendedprice"] * (1 - df["l_discount"]))
```

```
    .pipe(lambda df: df[rescols])
```

```
    .groupby(["l_orderkey", "o_orderdate", "o_shippriority"], as_index=False)
```

```
    .agg({"revenue": "sum"})[rescols]
```

```
    .sort_values(["revenue", "o_orderdate"], ascending=[False, True])
```

```
    .head(10)
```

```
)
```

Pandas Specific Optimization – Parameter Tuning

department-wise average salaries sorted in descending order

parameter tuning in pandas

```
res = (  
    employee.groupby("department")["salary"]  
        .mean()  
        .sort_values(ascending=False)  
)
```

```
res = (  
    employee.groupby("department", sort=False)["salary"]  
        .mean()  
        .sort_values(ascending=False)  
)
```

department	salary (USD)
IT	85,000
Admin	60,000
Finance	100,000
IT	81,000
Finance	95,000
Corporate	78,000
Sales	80,000

employee table

department	salary (USD)
IT	85,000
IT	81,000

department	salary (USD)
Admin	60,000

department	salary (USD)
Finance	100,000
Finance	95,000

department	salary (USD)
Corporate	78,000

department	salary (USD)
Sales	80,000

creating groups

department	salary (USD)
IT	83,000
Admin	60,000
Finance	97,500
Corporate	78,000
Sales	80,000

group-wise average-salary

department	salary (USD)
Admin	60,000
Corporate	78,000
Finance	97,500
IT	83,000
Sales	80,000

group-wise average-salary

sorted by "department"

department	salary (USD)
Finance	97,500
IT	83,000
Sales	80,000
Corporate	78,000
Admin	60,000

group-wise average-salary
sorted by "department"

```
df.groupby(["A", "B"])["C"]  
    .mean()  
    .sort_values(ascending=False)  
)
```

~50 sec

```
df.groupby(["A", "B"],  
    sort=False)["C"]  
    .mean()  
    .sort_values(ascending=False)
```


~30 sec

100M
samples with
high-
cardinality

Pandas Specific Optimization – Auto-selection of optimized method

Datetime Extractor

```
year = date.dt.strftime("%Y").astype(int)  
month = date.dt.strftime("%m").astype(int)  
day = date.dt.strftime("%d").astype(int)
```




Datetime Extractor

```
year = date.dt.year  
month = date.dt.month  
day = date.dt.day
```

Pandas Specific Optimization – Optimization on Index

```
sorted = df.sort_values("a").reset_index(drop=True)
```



```
sorted = df.sort_values("a", ignore_index=True)
```

Benchmark (1): DB-Benchmark

Database-like ops benchmark (<https://duckdblabs.github.io/db-benchmark>)

groupby join

0.5 GB 5 GB 50 GB

basic questions

rank-1

Input table: 1,000,000,000 rows x 9 columns (50 GB)

FireDucks	1.0.4	2024-09-10	15s
DuckDB	1.0.0	2024-07-04	25s
ClickHouse	24.5.1.1763	2024-06-07	28s
Polars	1.1.0	2024-07-09	47s
Datafusion	38.0.1	2024-06-07	56s
data.table	1.15.99	2024-06-07	88s
DataFrames.jl	1.6.1	2024-06-07	91s
InMemoryDataGrid	0.7.18	2023-10-17	218s
spark	3.5.1	2024-06-07	261s
R-arrow	16.1.0	2024-06-07	378s
collapse	2.0.14	2024-06-07	411s
(py)datatable	1.2.0a0	2024-06-07	1022s
dplyr	1.1.4	2024-06-07	1104s
pandas	2.2.2	2024-06-07	1126s
dask	2024.5.2	2024-06-07	out of memory
Modin		see README	pending

groupby join

0.5 GB 5 GB 50 GB

basic questions

rank-1

Input table: 100,000,000 rows x 7 columns (5 GB)

FireDucks	1.0.4	2024-09-10	7s
DuckDB	1.0.0	2024-07-04	9s
Polars	1.1.0	2024-07-08	9s
Datafusion	38.0.1	2024-06-07	15s
InMemoryDataGrid	0.7.18	2023-10-20	25s
ClickHouse	24.5.1.1763	2024-06-07	43s
data.table	1.15.99	2024-06-07	62s
collapse	2.0.14	2024-06-07	69s
DataFrames.jl	1.6.1	2024-06-07	77s
spark	3.5.1	2024-06-07	128s
dplyr	1.1.4	2024-06-07	214s
pandas	2.2.2	2024-06-07	244s
dask	2024.5.2	2024-06-07	635s
(py)datatable	1.2.0a0	2024-06-07	undefined exception
R-arrow	16.1.0	2024-06-07	out of memory
Modin		see README	pending


Benchmark (2): Speedup from pandas in TPC-H benchmark

FireDucks is ~345x faster than pandas at max

Server

Xeon Gold 5317 x2
(24 cores), 256GB

Speedup from pandas 2.2.2 (scale factor = 10)



Comparison of
DataFrame libraries
(average speedup)

FireDucks 50x


Polars 39x

Modin 0.9x

Benchmark (3): Speedup from pandas in TPCx-BB benchmark

ETL(Extract, Transform, Load) and ML Workflow

FireDucks speedup from pandas



- pandas-2.1.4
- fireducks-0.9.3

- CPU: Intel(R) Xeon(R) Gold 5317 CPU @ 3.00GHz x 2sockets (Total 48HW Threads)
- Main memory: 256GB

Resource on FireDucks

Web site (User guide, benchmark, blog)

<https://fireducks-dev.github.io/>



X(twitter) (Release information)

<https://x.com/fireducksdev>



Github (Issue report)

<https://github.com/fireducks-dev/fireducks>



slack Q/A, communication

https://join.slack.com/t/fireducks/shared_invite/zt-2j4lucmtj-IGR7AWIXO62Lu605pnBJ2



FireDucks

Compiler Accelerated DataFrame Library for Python with fully-compatible pandas API

Get Started

```
import fireducks.pandas as pd
```

News

[Release fireducks-0.12.4 \(Jul 09, 2024\)](#)

[Have you ever thought of speeding up your data analysis in pandas with a compiler?\(blog\) \(Jul 03, 2024\)](#)

[Evaluation result of Database-like ops benchmark with FireDucks is now available. \(Jun 18, 2024\)](#)



Accelerate pandas without any manual code changes

Do you have a pandas-based program that is slow? FireDucks can speed-up your programs without any manual code changes. You can accelerate your data analysis without worrying about slow performance due to single-threaded execution in pandas.

Let's go for a test drive!

<https://colab.research.google.com/drive/1qpej-X7CZsleOqKuhBg4kq-cbGuJf1Zp?usp=sharing>



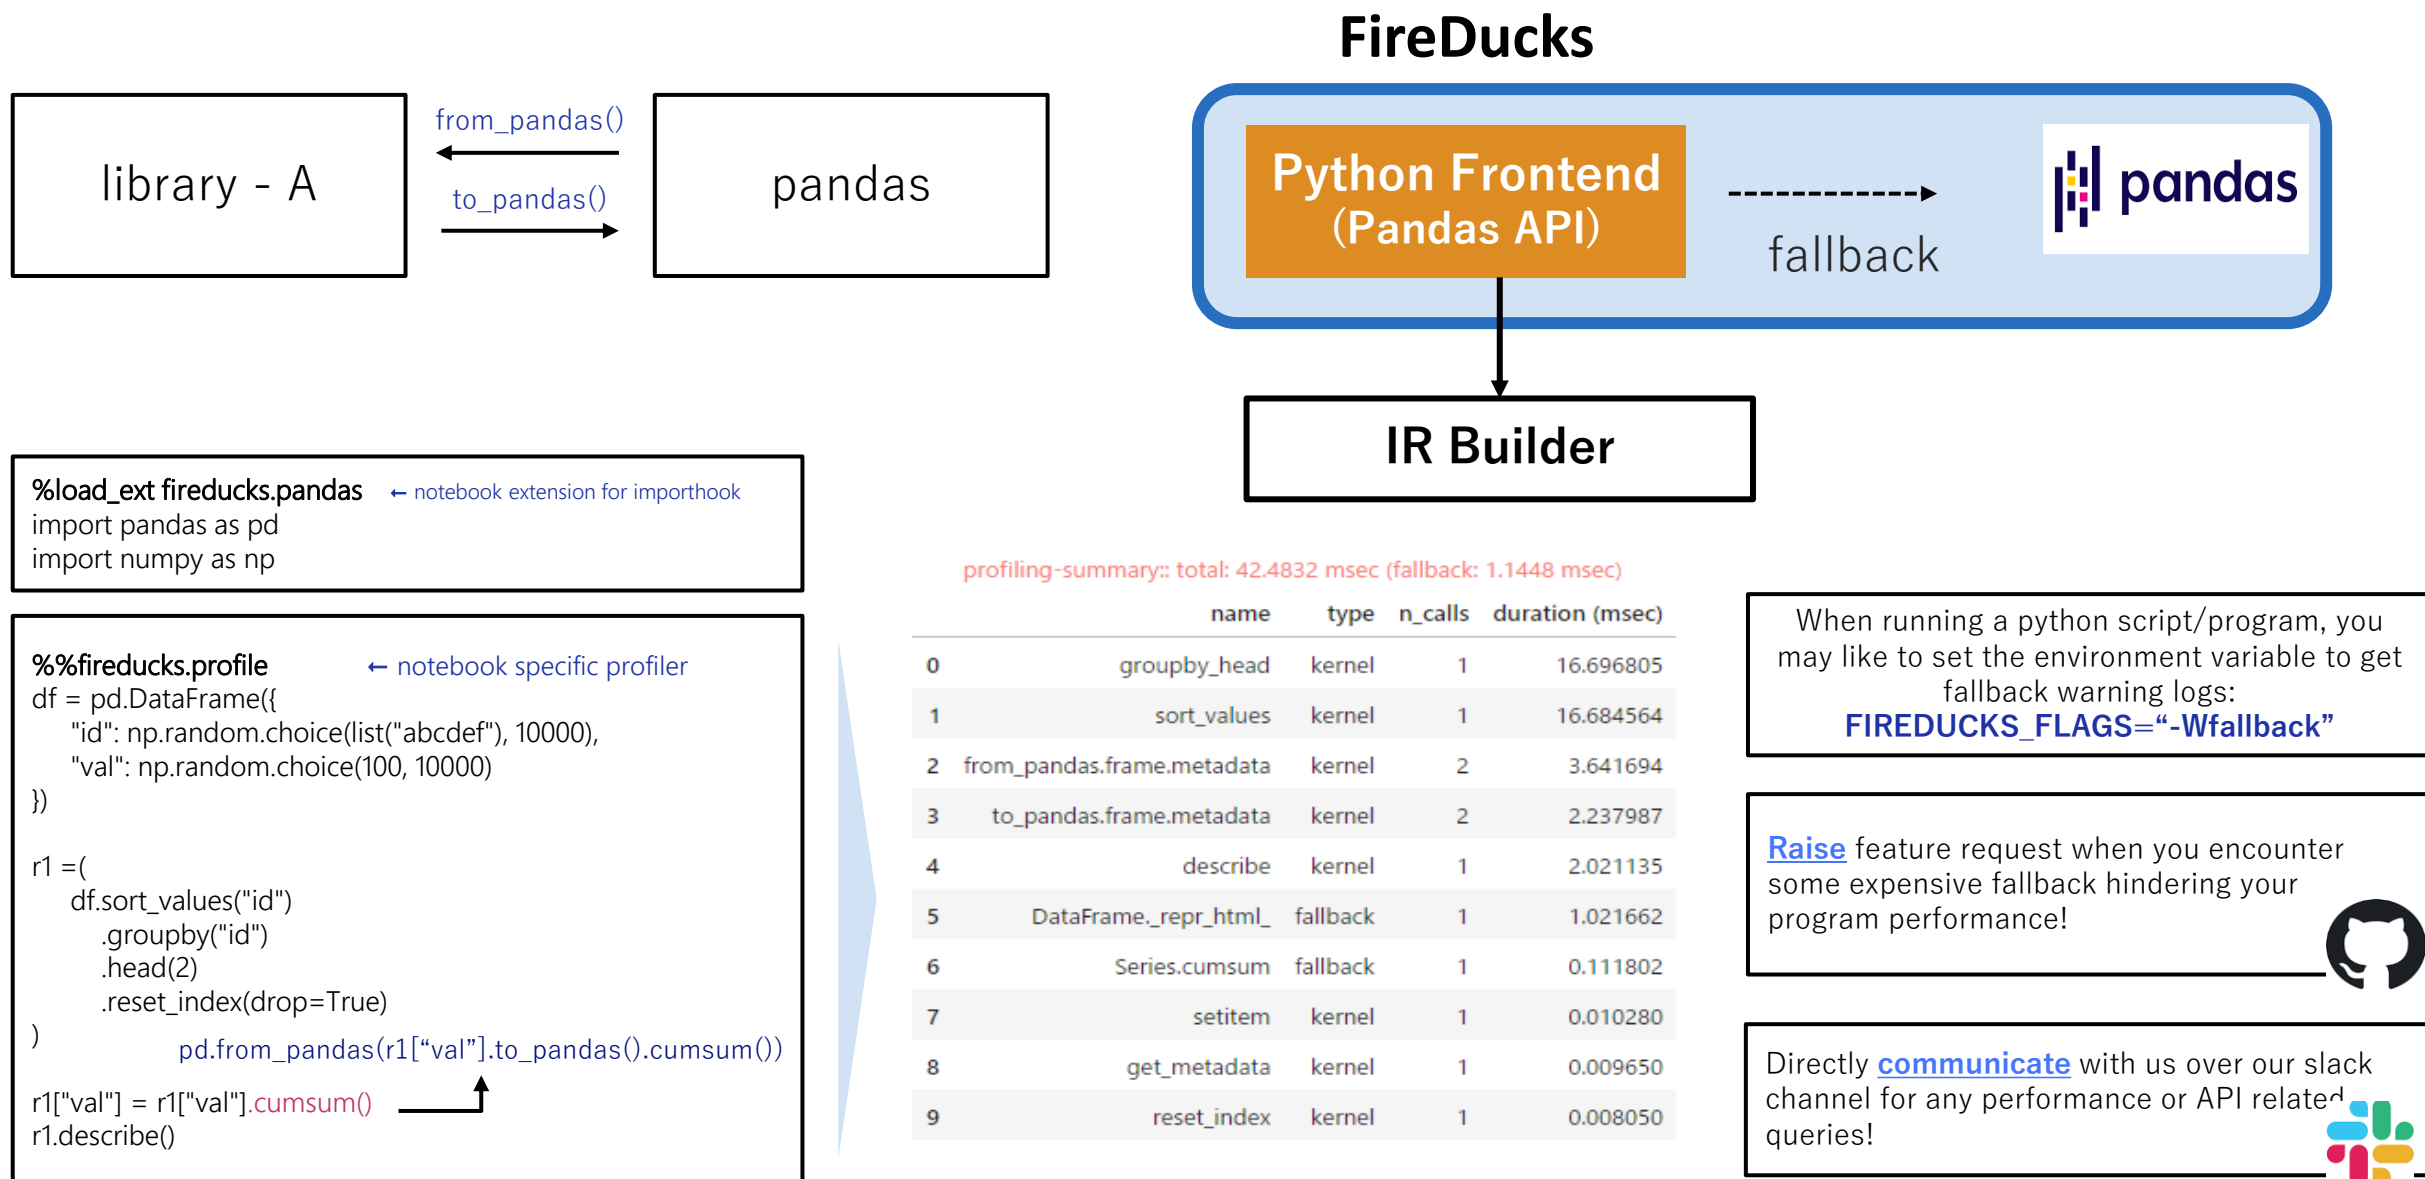
Thank You!

- ◆ Focus more on in-depth data exploration using “pandas”.
- ◆ Let the “FireDucks” take care of the optimization for you.
- ◆ Enjoy Green Computing!



Frequently Asked Questions

FAQ: Why FireDucks is highly compatible with pandas?



FAQ: How to evaluate Lazy Execution?

```
def foo(employee, country):  
    stime = time.time()  
    m = employee.merge(country, on="C_Code")  
    r = m[m["Gender"] == "Male"]  
    print(f"fireducks time: {time.time() - stime} sec")  
    return r
```

**fireducks time: 0.0000123
sec**

```
def foo(employee, country):  
    employee._evaluate()  
    country._evaluate()  
    stime = time.time()  
    m = employee.merge(country, on="C_Code")  
    r = m[m["Gender"] == "Male"]  
    r._evaluate()  
    print(f"fireducks time: {time.time() - stime} sec")  
    return r
```

**fireducks time: 0.02372143
sec**



IR Builder

```
create_data_op(...)  
merge_op(...)  
filter_op(...)
```

FIREDUCKS_FLAGS="--benchmark-mode"



Use this to disable lazy-execution mode when you do not want to make any changes in your existing application during performance evaluation.

FAQ: How to configure number of cores to be used?

OMP_NUM_THREADS=1



Use this to stop parallel execution, or configure this with the intended number of cores to be used



Alternatively, you can use the Linux taskset command to bind your program with specific CPU cores.



Orchestrating a brighter world

NECは、安全・安心・公平・効率という社会価値を創造し、
誰もが人間性を十分に発揮できる持続可能な社会の実現を目指します。

\Orchestrating a brighter world

NEC